

Last updated: May 21, 2019

System Design Template

(1) Use cases

- (2) Scenarios that will **not** be covered
- (3) Who will use
- (4) How many will use
- (5) Usage patterns

(2) ESTIMATIONS [5 min]

- (1) Throughput (QPS **for read and write** queries)
- (2) Latency expected **from** the system (**for read and write** queries)
- (3) **Read/Write** ratio
- (4) Traffic estimates
 - **Write** (QPS, Volume **of** data)
 - **Read** (QPS, Volume **of** data)
- (5) Storage estimates
- (6) Memory estimates
 - **If** we are **using** a cache, what **is** the kind **of** data we want **to** store **in** cache
 - How much RAM **and** how many machines **do** we need **for** us **to** achieve this ?
 - Amount **of** data you want **to** store **in** disk/ssd

(3) DESIGN GOALS [5 min]

- (1) Latency and Throughput requirements
- (2) Consistency vs Availability [Weak/strong/**eventual** => consistency | Failover/**replication** => availability]

(4) HIGH LEVEL DESIGN [5-10 min]

- (1) APIs **for Read/Write** scenarios **for** crucial components
- (2) **Database schema**
- (3) Basic algorithm
- (4) High **level** design **for Read/Write** scenario

(5) DEEP DIVE [15-20 min]

- (1) Scaling the algorithm
- (2) Scaling individual components:
 - > Availability, Consistency **and** Scale story **for** each component
 - > Consistency **and** availability patterns
- (3) Think about the following components, how they would fit **in and** how it would help
 - a) **DNS**
 - b) **CDN** [Push vs Pull]
 - c) **Load Balancers** [Active-Passive, Active-Active, Layer 4, Layer 7]
 - d) **Reverse Proxy**

>> Master-slave, Master-master, Federation, Sharding, Denormalization, SQL

Tuning

> NoSQL

>> Key-Value, Wide-Column, Graph, Document

Fast-lookups:

>>> RAM [Bounded size] => Redis, Memcached

>>> AP [Unbounded size] => Cassandra, RIAK, Voldemort

>>> CP [Unbounded size] => HBase, MongoDB, Couchbase, DynamoDB

g) Caches

> Client caching, CDN caching, Webserver caching, Database caching, Application caching, Cache @Query level, Cache @Object level

> Eviction policies:

>> Cache aside

>> Write through

>> Write behind

>> Refresh ahead

h) Asynchronism

> Message queues

> Task queues

> Back pressure

i) Communication

> TCP

> UDP

> REST

> RPC

(6) JUSTIFY [5 min]

- (1) Throughput of each layer
- (2) Latency caused between each layer
- (3) Overall latency justification

<https://www.gbmb.org/kb-to-mb>

| KB | MB | GB | TB | PB |
|----------|----------|----------|-----------|-----------|
| 10^3 | 10^6 | 10^9 | 10^{12} | 10^{15} |
| 2^{10} | 2^{20} | 2^{30} | 2^{40} | 2^{50} |

1 Million = 10^6

b. Datastore layer

We can divide our datastore layer into two:

1. Metadata database: We can use a relational database like MySQL or a Distributed Key-Value store like Dynamo or Cassandra.
2. Object storage: We can store our contents in an Object Storage like Amazon's S3. Whenever we feel like hitting our full capacity on content storage, we can easily increase it by adding more servers.

